

Störmer's Rule

Mathematical Derivation from First Principles

Cosmos GNC Library — Propulsion Repository

2026

Contents

1	Problem Statement	2
2	The Störmer Three-Point Recurrence	2
2.1	Derivation from Taylor's Theorem	2
2.2	The δ -Increment Formulation	3
3	Order Analysis: What “Order p” Means	3
3.1	The HNW Framework	3
3.2	Verification for the $k = 2$ Störmer Method	5
3.3	The Rounding-Error Trap: Why the δ -Form Is Mandatory	5
4	Störmer's Modified Midpoint Algorithm	6
5	Equivalence to Störmer-Verlet / Leapfrog	6
6	Even Error Series: Gragg's Theorem	7
7	Richardson Extrapolation Setup (Bulirsch-Stoer)	7
8	Analysis of the NR3 Implementation	8
8.1	Parameterization	8
8.2	Algorithm in NR3 variables	8
8.3	Cost comparison	9
9	Is Inheritance Necessary?	9
10	Harmonic Oscillator Test Case	9
11	Numerov's Method: The Implicit Upgrade	10
11.1	Notation Bridge: HNW Polynomials vs NR3/Our Variables	10
11.2	Derivation via the Order Condition	11
11.3	Order Verification	11
11.4	The Order-Barrier Context	11
11.5	PECECE Implementation	12

11.6 Cost Comparison	12
11.7 Starting Procedure	13
12 Astrodynamics Context: Cowell's Method, Encke's Method, and Perturbations	13
12.1 Cowell's Method (Direct Integration)	13
12.2 Encke's Method (Osculating Orbit)	14
12.3 Störmer/Numerov vs. Cowell for Orbital Mechanics	14
12.4 Common Perturbations in Astrodynamics	15
12.5 Application: Numerov Orbit Propagation Demo	15
13 Summary	16

1 Problem Statement

We consider a second-order ordinary differential equation of the form

$$y'' = f(x, y), \quad y(x_0) = y_0, \quad y'(x_0) = v_0, \quad (1)$$

where $y : \mathbb{R} \rightarrow \mathbb{R}^n$, $f : \mathbb{R} \times \mathbb{R}^n \rightarrow \mathbb{R}^n$, and crucially f does not depend on y' .

Remark 1 (Physical motivation). *Equation (1) describes any purely position-dependent force: gravitational n -body, Keplerian orbits, conservative mechanical systems, and the Störmer problem (charged particle in a magnetic field) that originally motivated the method. For the two-body orbital problem: $f(\mathbf{r}) = -\mu \mathbf{r}/|\mathbf{r}|^3$.*

Goal. Advance the solution from x_0 to $x_0 + H$ by subdividing H into M equal sub-steps of size $h = H/M$, in a way that:

1. exploits the absence of y' in f ,
2. produces an *even* error series in h (enabling Bulirsch-Stoer Richardson extrapolation), and
3. requires only $M/2$ new evaluations of f per column (half the cost of applying the standard modified midpoint rule to the equivalent first-order system).

2 The Störmer Three-Point Recurrence

2.1 Derivation from Taylor's Theorem

Write the Taylor expansions of y centered at $x_m = x_0 + mh$:

$$y(x_m + h) = y_m + hy'_m + \frac{h^2}{2}y''_m + \frac{h^3}{6}y'''_m + \frac{h^4}{24}y^{(4)}_m + \mathcal{O}(h^5), \quad (2)$$

$$y(x_m - h) = y_m - hy'_m + \frac{h^2}{2}y''_m - \frac{h^3}{6}y'''_m + \frac{h^4}{24}y^{(4)}_m + \mathcal{O}(h^5). \quad (3)$$

Adding (2) and (3) cancels all odd powers:

$$y(x_m + h) + y(x_m - h) = 2y_m + h^2 y_m'' + \frac{h^4}{12} y_m^{(4)} + \mathcal{O}(h^6). \quad (4)$$

Substituting $y_m'' = f(x_m, y_m)$ and rearranging:

$$\boxed{y_{m+1} = 2y_m - y_{m-1} + h^2 f(x_m, y_m) + \mathcal{O}(h^4).} \quad (5)$$

This is the **Störmer three-point formula**: given two successive solution values y_{m-1} and y_m , the next value y_{m+1} requires a single evaluation of f and is accurate to $\mathcal{O}(h^4)$ locally.

Remark 2 (No velocity in the recurrence). *Equation (5) uses y_{m-1} , y_m , and $f(x_m, y_m)$ only. The velocity y' has been eliminated algebraically. The price is the need for two starting values.*

2.2 The δ -Increment Formulation

Define the position increment $\delta_m := \theta_{m+1} - \theta_m$, the change in the numerical solution over one sub-step. From (5):

$$\begin{aligned} \delta_m &= y_{m+1} - y_m \\ &= (2y_m - y_{m-1} + h^2 f_m) - y_m \\ &= (y_m - y_{m-1}) + h^2 f_m \\ &= \delta_{m-1} + h^2 f_m. \end{aligned} \quad (6)$$

This telescoping relation propagates δ forward with only $h^2 f_m$ added at each step. The velocity information is encoded in the initial δ_0 .

Starting value from Taylor: The value $y_1 = y(x_0 + h)$ from (2) with $m = 0$:

$$y_1 = y_0 + h v_0 + \frac{h^2}{2} f_0 + \mathcal{O}(h^3), \quad (7)$$

so

$$\boxed{\delta_0 = h v_0 + \frac{h^2}{2} f_0.} \quad (8)$$

3 Order Analysis: What “Order p ” Means

The three-point recurrence (5) was derived from Taylor’s theorem. We now ask precisely: *how fast does the global error shrink as $h \rightarrow 0$?* The answer, “order p ”, refers entirely to powers of h — it is **not** the degree of a polynomial fitted to f .

3.1 The HNW Framework

Following Hairer, Nørsett & Wanner [1] (henceforth **HNW**), write the general second-order multistep method for $y'' = f(x, y)$ as

$$\sum_{i=0}^k \alpha_i y_{n+i} = h^2 \sum_{i=0}^k \beta_i f(x_{n+i}, y_{n+i}), \quad \text{HNW (III.10, Eq. 10.19)} \quad (9)$$

with scalar sequences $\{\alpha_i\}$ and $\{\beta_i\}$, $\alpha_k \neq 0$. Define the two **characteristic polynomials**

$$\varrho(\zeta) = \sum_{i=0}^k \alpha_i \zeta^i, \quad \sigma(\zeta) = \sum_{i=0}^k \beta_i \zeta^i. \quad (10)$$

These encode the entire method: ϱ governs stability and σ governs accuracy.

The residual when the *exact* solution is substituted into the method is measured by the **linear difference operator** (HNW, Eq. 10.23, p. 467):

$$\mathcal{L}[y, x, h] = \varrho(E) y(x) - h^2 \sigma(E) y''(x) = \sum_{i=0}^k \left[\alpha_i y(x + ih) - h^2 \beta_i y''(x + ih) \right], \quad (11)$$

where E is the shift operator $E y(x) = y(x + h)$. Along the exact solution $y''(x + ih) = f(x + ih, y(x + ih))$, so $\mathcal{L}$ is a pure function of the smooth exact solution y and the step h .

Theorem 1 (HNW Definition 10.2 and Theorem 10.3, p. 467–468). *Method (9) is **consistent of order** p if*

$$\mathcal{L}[y, x, h] = \mathcal{O}(h^{p+2}) \quad \text{for every sufficiently smooth } y. \quad (12)$$

Three equivalent conditions (HNW Theorem 10.3, p. 468) are:

(i) Moment form. $\sum_i \alpha_i = 0$, $\sum_i i \alpha_i = 0$, and for $q = 2, \dots, p+1$:

$$\sum_{i=0}^k \alpha_i i^q = q(q-1) \sum_{i=0}^k \beta_i i^{q-2}. \quad (13)$$

(ii) Generating-polynomial form.

$$\varrho(e^h) - h^2 \sigma(e^h) = \mathcal{O}(h^{p+2}) \quad \text{as } h \rightarrow 0. \quad (14)$$

(iii) Root form near $\zeta = 1$.

$$\frac{\varrho(\zeta)}{(\log \zeta)^2} - \sigma(\zeta) = \mathcal{O}((\zeta - 1)^p) \quad \text{as } \zeta \rightarrow 1. \quad (15)$$

Why h^{p+2} , not h^p ? The method equates the h^2 -scaled forcing $h^2 f$ to the finite-difference $\sum \alpha_i y_{n+i}$, so the entire equation is naturally $\mathcal{O}(h^2)$. The residual $\mathcal{L}$ is therefore already premultiplied by h^2 ; order p means the *extra* smallness beyond h^2 is h^p more — i.e., h^{p+2} total. The global error, obtained by accumulating $\mathcal{O}(h^{p+2})$ local residuals over $\mathcal{O}(1/h)$ steps and dividing once by the implicit h^2 factor, is $\mathcal{O}(h^p)$ — matching Definition 10.2. This is confirmed by:

Theorem 2 (HNW Theorem 10.6, p. 468 — Global Convergence). *If method (9) is stable and of order p , and the starting values satisfy $y(x_j) - y_j = \mathcal{O}(h^{p+1})$ for $j = 0, \dots, k-1$, then*

$$|y(x_n) - y_n| \leq C h^p \quad \text{for all } 0 \leq hn \leq \text{const}. \quad (16)$$

3.2 Verification for the $k = 2$ Störmer Method

For the three-point formula (5):

$$\alpha_0 = 1, \alpha_1 = -2, \alpha_2 = 1; \quad \beta_0 = 0, \beta_1 = 1, \beta_2 = 0. \quad (17)$$

Hence $\varrho(\zeta) = \zeta^2 - 2\zeta + 1 = (\zeta - 1)^2$ and $\sigma(\zeta) = 1$ (the center-point quadrature rule for f).

Apply the generating-polynomial test (14). Expanding $\varrho(e^h)$ term by term:

$$\varrho(e^h) = e^{2h} - 2e^h + 1 = h^2 + \frac{h^4}{12} + \frac{h^6}{360} + \cdots \quad (18)$$

(all odd powers cancel because $e^{2h} - 2e^h + 1$ is symmetric under $h \mapsto -h$ after recentering). With $\sigma(\zeta) = 1$:

$$\varrho(e^h) - h^2\sigma(e^h) = \underbrace{h^2}_{\text{cancel}} + \frac{h^4}{12} + \frac{h^6}{360} + \cdots - h^2 = \frac{h^4}{12} + \frac{h^6}{360} + \cdots = \mathcal{O}(h^4). \quad (19)$$

Comparing with condition (12): $h^{p+2} = h^4 \Rightarrow \boxed{p = 2}$. The method is **globally second-order accurate** in h .

The leading residual coefficient $\frac{1}{12}$ is the **error constant** $C = \sigma_2 = \frac{1}{12}$ (HNW Exercise 5, §III.10). It governs the asymptotic error via (HNW Eq. 10.36–10.37, p. 471):

$$\kappa''(x) = \frac{\partial f}{\partial y}(x, y(x)) \kappa(x) - \frac{1}{12} y^{(4)}(x), \quad (20)$$

where $\kappa(x)$ satisfies $y_n - y(x_n) = \kappa(x_n) h^2 + \mathcal{O}(h^3)$. For orbital mechanics, $y^{(4)}$ is the fourth derivative of the position trajectory; a smooth orbit (low eccentricity) keeps $y^{(4)}$ small, making the error constant favorable.

3.3 The Rounding-Error Trap: Why the δ -Form Is Mandatory

The polynomial $\varrho(\zeta) = (\zeta - 1)^2$ has a **double root at** $\zeta = 1$. HNW Definition 10.1 (p. 467) allows roots on the unit circle with multiplicity at most 2 — our method sits exactly on this boundary.

What this means in practice (HNW §III.10, pp. 472–473, Fig. 10.3): a *direct* implementation that stores $y_0, y_1, \dots$ and applies the recurrence $y_{n+1} = 2y_n - y_{n-1} + h^2 f_n$ accumulates rounding errors that grow without bound as $h \rightarrow 0$. The double root causes the companion matrix of the difference equation to be no longer power-bounded.

The **fix** is the δ -increment form (HNW Eq. 10.38', p. 472):

$$\delta_n - \delta_{n-1} = h^2 f_n, \quad y_{n+1} - y_n = \delta_n. \quad (21)$$

This splits $\varrho(\zeta)$ as $(\zeta - 1)(\zeta - 1)$ and introduces $\delta_n = y_{n+1} - y_n$ as a new variable. The resulting companion matrix is power-bounded (HNW Eq. 10.30), and rounding errors remain bounded. This is exactly the formulation used in `StormerStep::step()`.

A second pitfall (HNW p. 473, note b): when computing *starting values* $\delta_0, \delta_1, \dots$, do **not** form the differences $(y_1 - y_0)/h$, $(y_2 - y_1)/h$ from separately computed RK values — that difference is numerically unstable. Instead, accumulate the increments directly from the sub-steps of the RK starting procedure, just as `StormerStep` does internally.

4 Störmer's Modified Midpoint Algorithm

Notation. We introduce θ_m to denote the **numerical approximation** to the true solution $y(x_m)$ at sub-step m , where $x_m = x_0 + mh$. The exact solution is $y(x_m)$; its approximation is θ_m . The algorithm computes a sequence of approximations:

$$\theta_0 \approx y(x_0), \quad \theta_1 \approx y(x_1), \quad \dots, \quad \theta_M \approx y(x_0 + H).$$

The velocity approximation at the final point is denoted $z_M \approx y'(x_0 + H)$.

Why θ instead of y ? In numerical analysis, it is standard to distinguish the *exact* solution $y(x)$ (unknown, continuous) from the *computed* approximation θ_m (known, discrete). The error at step m is $y(x_m) - \theta_m$. This distinction becomes crucial when analyzing convergence and error propagation.

The complete algorithm for M sub-steps of size $h = H/M$ is:

1. **Initialize.**

$$\theta_0 = y_0, \tag{22}$$

$$\delta_0 = h v_0 + \frac{h^2}{2} f(x_0, \theta_0). \tag{23}$$

Set $\theta_1 = \theta_0 + \delta_0$.

2. **Sub-step loop** for $m = 1, \dots, M - 1$:

$$f_m = f(x_0 + mh, \theta_m), \tag{24}$$

$$\delta_m = \delta_{m-1} + h^2 f_m, \tag{25}$$

$$\theta_{m+1} = \theta_m + \delta_m. \tag{26}$$

3. **Final derivative** (velocity estimate at $x_0 + H$):

$$z_M = \frac{\delta_{M-1}}{h} + \frac{h}{2} f(x_0 + H, \theta_M). \tag{27}$$

The output is $(\theta_M, z_M) \approx (y(x_0 + H), y'(x_0 + H))$, the numerical approximations to the true position and velocity at $x_0 + H$, with local truncation error $\mathcal{O}(h^2)$.

Total function evaluations of f : one per sub-step $m = 0, \dots, M$, i.e. $M + 1$ total, but f_0 is supplied from the previous step, so M *new evaluations*.

5 Equivalence to Störmer-Verlet / Leapfrog

Define the half-integer velocity $v_{m+1/2} := \delta_m/h$. Rewriting the algorithm in terms of v :

$$v_{1/2} = \frac{\delta_0}{h} = v_0 + \frac{h}{2} f_0 \quad (\text{half-kick}), \tag{28}$$

$$\theta_{m+1} = \theta_m + h v_{m+1/2} \quad (\text{drift}), \tag{29}$$

$$v_{m+3/2} = v_{m+1/2} + h f_{m+1} \quad (\text{full kick, } m = 0, \dots, M - 2), \tag{30}$$

and the final velocity:

$$z_M = v_{M-1/2} + \frac{h}{2} f_M. \quad (31)$$

Equation (31) can be read as: take the half-integer velocity $v_{M-1/2}$, and kick it forward by a half step to return to integer time M . Equivalently:

$$z_M = \frac{v_{M-1/2} + v_{M+1/2}}{2}, \quad \text{where } v_{M+1/2} = v_{M-1/2} + h f_M. \quad (32)$$

So z_M is the arithmetic mean of the bracketing half-integer velocities, a second-order approximation to $y'(x_0 + H)$.

Remark 3 (Symplecticity). *The drift-kick-drift structure (28)–(30) is the Störmer-Verlet integrator, which is symplectic for Hamiltonian systems $H = \frac{1}{2}|v|^2 + V(y)$ with $f = -\nabla V$. It preserves a modified Hamiltonian exactly to all orders. Importantly, it does NOT suffer from secular energy drift.*

6 Even Error Series: Gragg’s Theorem

Theorem 3 (Gragg 1965). *Let $y(x)$ be sufficiently smooth. For Störmer’s modified midpoint rule applied to $y'' = f(x, y)$ with M sub-steps of $h = H/M$ and the starting/ending conventions (8)–(27), the global errors admit expansions in even powers of h only:*

$$y(x_0 + H) - \theta_M = a_2(H) h^2 + a_4(H) h^4 + a_6(H) h^6 + \dots, \quad (33)$$

$$y'(x_0 + H) - z_M = b_2(H) h^2 + b_4(H) h^4 + b_6(H) h^6 + \dots, \quad (34)$$

where the coefficients $a_{2k}(H)$, $b_{2k}(H)$ are independent of M .

Sketch of proof (time-reversibility argument). The Störmer recurrence (22)–(27) is *time-reversible*: reversing $h \rightarrow -h$ and the sequence of θ_m recovers the original sequence. A time-reversible method cannot accumulate odd powers of h in its global error, because odd-power terms change sign under $h \rightarrow -h$ while the exact solution does not. More formally, one expands $\theta_m(h)$ as a power series in h and shows that all odd-power terms cancel identically (see Hairer, Nørsett & Wanner, Vol. I, §II.8 for the complete argument).

Corollary 1 (Richardson extrapolation efficiency). *Because the error series (33) contains only $h^2, h^4, \dots$, a single Richardson extrapolation step using results at h and $h/2$ eliminates the h^2 term and yields $\mathcal{O}(h^4)$ accuracy. In general, k extrapolation steps yield $\mathcal{O}(h^{2k+2})$ accuracy from $\mathcal{O}(h^2)$ base integrations. This is the foundation of the Bulirsch-Stoer method.*

7 Richardson Extrapolation Setup (Bulirsch-Stoer)

Evaluate the modified midpoint rule with sub-step counts $M_1 < M_2 < M_3 < \dots$ (the Bulirsch-Stoer sequence: $\{2, 4, 6, 8, 12, 16, \dots\}$), giving approximations $T_{j,0} := \theta_{M_j}(H)$.

Define the extrapolation table recursively:

$$T_{j,k} = T_{j,k-1} + \frac{T_{j,k-1} - T_{j-1,k-1}}{(M_j/M_{j-k})^2 - 1}, \quad k = 1, 2, \dots, j. \quad (35)$$

Each extrapolation step eliminates the next order in h^2 . The diagonal $T_{k,k}$ converges to $y(x_0 + H)$ as k increases.

Remark 4 (Computational cost vs. accuracy). *For the first-order system $[y', f(x, y)]^T$ of dimension $2n$, the modified midpoint rule needs M evaluations of f per column. Störmer's rule needs only $M/2$ evaluations per column (the NR3 code uses $n_{\text{eff}} = M/2$ leapfrog steps of $h_{\text{eff}} = 2h = 2H/M$, as shown below) for the same total interval H . Since f can be expensive (e.g., gravitational force with many bodies), this factor of two is significant.*

8 Analysis of the NR3 Implementation

8.1 Parameterization

NR3's `StepperStoerm::dy()` uses:

- `nstep` = M (from Bulirsch-Stoer sequence $\{2, 4, 6, 8, \dots\}$).
- `h` = H/M (unit sub-step).
- `h2` = $2h = 2H/M$ (actual loop step, double the unit sub-step).
- `nstp2` = $M/2$ (number of loop iterations).

8.2 Algorithm in NR3 variables

Initialization:

$$\begin{aligned} \text{ytemp}[i] &\leftarrow y[i] & \theta_0 &= y_0, \\ \text{ytemp}[n+i] &\leftarrow y'[i] + h \cdot f_0 & v^* &= v_0 + h f_0 = v_{1/2} \cdot \underbrace{2}_{h_{\text{eff}}/h}. \end{aligned}$$

In other words, the stored “velocity” `ytemp` $[n+i]$ is the half-integer velocity $v_{1/2} = v_0 + (h/2) f_0$ scaled by 2: `ytemp` $[n+i] = 2 v_{1/2}$. Then the drift step below with factor $h_{\text{eff}} = 2h$ is:

$$\theta_1 = \theta_0 + h_{\text{eff}} \cdot \text{ytemp}[n+i]/2 = \theta_0 + h_{\text{eff}} v_{1/2} = \theta_0 + 2h (v_0 + \frac{h}{2} f_0) = \theta_0 + 2h v_0 + h^2 f_0.$$

Alternatively, viewing the NR3 scheme directly as Störmer-Verlet with step $h_{\text{eff}} = 2h$:

$$\text{Loop } m = 1, \dots, n_{\text{eff}} : \quad \begin{cases} \theta_m = \theta_{m-1} + h_{\text{eff}} v_{m-1/2}, \\ v_{m+1/2} = v_{m-1/2} + h_{\text{eff}} f(x_0 + m h_{\text{eff}}, \theta_m), \end{cases} \quad m < n_{\text{eff}}. \quad (36)$$

Final velocity: $z = v_{n_{\text{eff}}-1/2} + (h_{\text{eff}}/2) f(x_0 + H, \theta_{n_{\text{eff}}})$.

8.3 Cost comparison

Method	State dimension	Evals of f per column	Error order
Midpoint (first-order 2n-system)	$2n$	M	$\mathcal{O}(h^2)$
Störmer (NR3)	n	$M/2$	$\mathcal{O}(h_{\text{eff}}^2)$

Both methods embed in a Bulirsch-Stoer extrapolation with an even-power error series and converge to the same accuracy. Störmer uses half the evaluations of f at the cost of doubling the effective sub-step size, which the extrapolation corrects.

9 Is Inheritance Necessary?

NR3's design uses `struct StepperStoerm : StepperBS<D> solely` for code reuse of the Bulirsch-Stoer extrapolation machinery (odeint driver, column sequence, Richardson table, stepsize control). The method `dy()` is entirely independent of the base class algorithms; it overrides only the *sub-step integration* while reusing everything else.

From first principles, the minimal reusable components are:

1. **StormerStep**: executes n_{eff} leapfrog sub-steps over a given interval H . No base class required. Template on `DerivativeType` and `ContainerT`.
2. **Error estimation**: step-doubling (compare n_{eff} steps vs. $2n_{\text{eff}}$ steps) or Richardson extrapolation as a separate, composable component.
3. **Step-size control**: compose with the existing `ComputePISetpSize / PISetpSizeControl` classes already present in the library.

Composition over inheritance, exactly consistent with how `CalculateNewYAndError` and `IntegrateWithPIControl` were designed.

10 Harmonic Oscillator Test Case

The standard test for Störmer's rule is the harmonic oscillator:

$$y'' = -\omega^2 y, \quad y(0) = 1, \quad y'(0) = 0, \quad \omega = 1. \quad (37)$$

Exact solution: $y(t) = \cos t$, $y'(t) = -\sin t$, with conserved energy:

$$E(t) = \frac{1}{2} (y'(t))^2 + \frac{1}{2} (y(t))^2 = \frac{1}{2}. \quad (38)$$

Störmer-Verlet (symplectic) conserves a modified energy exactly, so the energy error remains bounded for all time (no secular drift), unlike non-symplectic methods.

Expected convergence. For a single step of size $H = 0.5$ with n_{eff} leapfrog sub-steps of $h_{\text{eff}} = H/n_{\text{eff}}$:

$$|y(H) - \theta_M| \approx \frac{H^3}{24} |y'''(0)| \cdot h_{\text{eff}}^2 = \frac{H^3 \omega^3}{24} \cdot h_{\text{eff}}^2 = \mathcal{O}(h_{\text{eff}}^2).$$

Doubling n_{eff} (halving h_{eff}) should reduce the error by a factor of 4. This is the convergence test in the unit tests.

11 Numerov’s Method: The Implicit Upgrade

Section 3 showed that the three-point stencil (5) achieves order 2 because $\varrho(e^h) - h^2\sigma(e^h) = \mathcal{O}(h^4)$ with $\sigma(\zeta) = 1$. The same stencil on the left-hand side, but a richer choice of σ , can reach $\mathcal{O}(h^6)$ — i.e. order 4 — with *no extra history*.

11.1 Notation Bridge: HNW Polynomials vs NR3/Our Variables

The analysis in §3 introduced Hairer–Nørsett–Wanner (HNW) notation. Here is a translation table bridging their polynomial formalism with the discrete variables used earlier in this document and in NR3:

Symbol	Meaning	Where it appears
$\varrho(\zeta)$	First characteristic polynomial $\sum_{i=0}^k \alpha_i \zeta^i$	HNW Eq. 10.20, p. 467 (stability)
$\sigma(\zeta)$	Second characteristic polynomial $\sum_{i=0}^k \beta_i \zeta^i$	HNW Eq. 10.20, p. 467 (accuracy)
ζ	Formal variable; $E = \zeta$ is the shift operator: $Ey_n = y_{n+1}$	HNW §III.1
α_i	Coefficients of LHS stencil (position weights)	Eq. (9)
β_i	Coefficients of RHS quadrature (force weights)	Eq. (9)
k	Step number (history length k) Method uses $y_n, \dots, y_{n+k}$	HNW §III.10
h	Step size	Same throughout
H	Total interval (macro-step)	§4
y_n	Exact solution at $x_n = x_0 + nh$	Continuous $y(x_n)$
θ_n	Computed approximation to y_n	§4
f_n	Force $f(x_n, y_n)$ or $f(x_n, \theta_n)$	§4
δ_n	Position increment $\theta_{n+1} - \theta_n$	Eq. (6)
$v_{1/2}$	Half-integer velocity (leapfrog)	Eq. (28)

Connection to our earlier notation. In §2–4 we wrote the method as

$$\theta_{m+1} - 2\theta_m + \theta_{m-1} = h^2 f_m.$$

In HNW polynomial form this is

$$\varrho(\zeta) = \zeta^2 - 2\zeta + 1 = (\zeta - 1)^2, \quad \sigma(\zeta) = 1.$$

The coefficients are $(\alpha_0, \alpha_1, \alpha_2) = (1, -2, 1)$ and $(\beta_0, \beta_1, \beta_2) = (0, 1, 0)$. The shift operator E acts as $E\theta_m = \theta_{m+1}$, so the compact operator form is $\varrho(E)\theta_m = h^2\sigma(E)f_m$.

NR3 notation. NR3 uses ‘nstep’ for the number of sub-steps M , ‘h’ for H/M , ‘h2’ for $2h$, and ‘nstp2’ for $M/2$. Their ‘ytemp[n+i]’ stores the half-integer velocity $v_{m+1/2}$ scaled by 2 (see §??).

11.2 Derivation via the Order Condition

Keep $\varrho(\zeta) = \zeta^2 - 2\zeta + 1$ and seek $\sigma(\zeta) = \beta_0 + \beta_1\zeta + \beta_2\zeta^2$ (symmetric: $\beta_0 = \beta_2$) such that $\varrho(e^h) - h^2\sigma(e^h) = \mathcal{O}(h^6)$.

From (19), $\varrho(e^h) = h^2 + \frac{h^4}{12} + \frac{h^6}{360} + \dots$. We need $h^2\sigma(e^h)$ to match through the h^5 term. Expanding with $\beta_0 = \beta_2 \equiv \beta$, $\beta_1 \equiv \gamma$:

$$h^2\sigma(e^h) = h^2(\beta e^{2h} + \gamma e^h + \beta) = h^2[(2\beta + \gamma) + 2h\beta + h^2(2\beta + \frac{\gamma}{2}) + \dots].$$

Matching through h^5 against $\varrho(e^h)$ forces $2\beta + \gamma = 1$ (the h^2 term), and solving the system for the h^4 match gives $\beta = \frac{1}{12}$, $\gamma = \frac{10}{12}$. Hence:

$$\sigma(\zeta) = \frac{\zeta^2 + 10\zeta + 1}{12}. \quad (39)$$

The resulting method is (HNW Eq. 10.12, p. 464; B. Numerov 1924):

$$\boxed{y_{n+1} - 2y_n + y_{n-1} = \frac{h^2}{12}(f_{n+1} + 10f_n + f_{n-1})}, \quad (40)$$

called **Numerov's method**. The only difference from (5) is the f_{n+1} term on the right: one additional force evaluation at the new position.

11.3 Order Verification

Apply the generating-polynomial test (14):

$$\begin{aligned} h^2\sigma(e^h) &= \frac{h^2}{12}(e^{2h} + 10e^h + 1) \\ &= h^2 + \frac{h^4}{12} + \frac{h^6}{144} + \dots \end{aligned} \quad (41)$$

and recall $\varrho(e^h) = h^2 + \frac{h^4}{12} + \frac{h^6}{360} + \dots$. Subtracting:

$$\varrho(e^h) - h^2\sigma(e^h) = \left(\frac{1}{360} - \frac{1}{144}\right)h^6 + \mathcal{O}(h^8) = -\frac{1}{240}h^6 + \mathcal{O}(h^8) = \mathcal{O}(h^6). \quad (42)$$

Comparing with (12): $h^{p+2} = h^6 \Rightarrow \boxed{p = 4}$. **Numerov's method is globally fourth-order accurate.**

The error constant is $C = \sigma_3^* = 0$ (Table 10.2, HNW p. 464) — the leading error is $-h^6/240 \cdot y^{(6)}$ per step, contributing:

$$\kappa''(x) = \frac{\partial f}{\partial y} \kappa - \frac{1}{240} y^{(6)}(x) \quad (\text{global error: } y_n - y(x_n) = \kappa(x_n) h^4 + \mathcal{O}(h^5)). \quad (43)$$

11.4 The Order-Barrier Context

HNW Theorem 10.4 (p. 468) states that for a stable k -step Störmer method:

$$p \leq k + 2 \text{ if } k \text{ is even,} \quad p \leq k + 1 \text{ if } k \text{ is odd.} \quad (44)$$

For $k = 2$ (our stencil): $p_{\max} = 4$. Numerov saturates this bound — it is the *highest-order stable method* on the three-point stencil. The explicit method (order 2) leaves two orders on the table.

11.5 PECECE Implementation

Numerov's method is implicit: y_{n+1} appears on both sides via $f_{n+1} = f(x_{n+1}, y_{n+1})$. For nonlinear f (e.g. gravitational force), we solve using **PECECE** (Predict–Evaluate–Correct–Evaluate–Correct–Evaluate):

1. **Predict.** Use the explicit Störmer formula as first guess:

$$y_{n+1}^{(P)} = 2y_n - y_{n-1} + h^2 f_n. \quad (45)$$

2. **Evaluate.** $f_{n+1}^{(P)} = f(x_{n+1}, y_{n+1}^{(P)})$.

3. **Correct.**

$$y_{n+1}^{(1)} = 2y_n - y_{n-1} + \frac{h^2}{12} (f_{n+1}^{(P)} + 10f_n + f_{n-1}). \quad (46)$$

4. **Evaluate.** $f_{n+1}^{(1)} = f(x_{n+1}, y_{n+1}^{(1)})$.

5. **Correct again.**

$$y_{n+1} = 2y_n - y_{n-1} + \frac{h^2}{12} (f_{n+1}^{(1)} + 10f_n + f_{n-1}). \quad (47)$$

6. **Evaluate.** $f_{n+1} = f(x_{n+1}, y_{n+1})$ (stored for next step).

Velocity. Numerov delivers positions. Velocity at integer time is recovered via the leapfrog formula (order 2 in h):

$$v_{n+1} = \frac{y_{n+1} - y_n}{h} + \frac{h}{2} f_{n+1}. \quad (48)$$

Position accuracy is $\mathcal{O}(h^4)$; velocity accuracy is $\mathcal{O}(h^2)$. For applications where velocity precision matters as much as position (e.g. orbit determination), the velocity can be improved via central differences: $v_n \approx (y_{n+1} - y_{n-1})/(2h) + \mathcal{O}(h^2)$ — same order.

Order of PECECE. With an order-2 predictor and order-4 corrector, PECE alone is order 3: the predictor error $\mathcal{O}(h^2)$ contaminates $f_{n+1}^{(P)}$ by $\mathcal{O}(h^2)$, adding $\frac{h^2}{12} \cdot \mathcal{O}(h^2) = \mathcal{O}(h^4)$ per step to the corrector residual, giving $\mathcal{O}(h^3)$ global error. The second correction (PECECE) reduces the contamination to $\mathcal{O}(h^6)$ per step — below the corrector's own leading error of $\mathcal{O}(h^6)$ — so the full order-4 accuracy is recovered.

11.6 Cost Comparison

Method	Order p	f evals/step	Error constant C
Explicit Störmer ($k = 2$)	2	1	$\sigma_2 = \frac{1}{12}$
Numerov PECECE ($k = 2$)	4	3	$-\frac{1}{240}$

Halving h reduces error by $2^2 = 4$ for explicit Störmer vs. $2^4 = 16$ for Numerov — a $4\times$ gain per step halving. The cost of 3 f -evals vs. 1 is recovered after the first tolerance tightening: to reach the same accuracy, Numerov needs $\approx 2.5\times$ fewer steps than explicit Störmer (since $4^{1/4}/2^{1/2} \approx 1.41$ ratio in step size for equal error).

11.7 Starting Procedure

Numerov is a two-step method: it needs y_0 and y_1 to start. Given initial conditions (y_0, v_0) , one accurate starting step is required. Theorem 2 demands $y(x_1) - y_1 = \mathcal{O}(h^5)$ for the method to retain its full order-4 convergence. Options:

- **Taylor expansion:** $y_1 = y_0 + hv_0 + \frac{h^2}{2}f_0 + \frac{h^3}{6}f'_0 + \frac{h^4}{24}f''_0$ (requires derivative information, gives $\mathcal{O}(h^5)$, satisfies requirement).
- **Runge-Kutta-Nyström order 4:** one step of a 4th-order Nyström method (HNW Table 14.2, p. 285) gives $\mathcal{O}(h^5)$ accuracy, at the cost of 3–4 f -evaluations for the single startup step.
- **Two steps of explicit Störmer** (order 2): gives only $\mathcal{O}(h^3)$, limiting global convergence to $\mathcal{O}(h^2)$. Sufficient only for testing.

`NumerovStep::startup()` uses one step of the leapfrog (Störmer-Verlet) with n_{eff} sub-steps as a high-accuracy starter: for large n_{eff} , this approximates the exact solution to whatever precision is needed. In production, supply a 4th-order starter.

12 Astrodynamics Context: Cowell’s Method, Encke’s Method, and Perturbations

In orbital mechanics, the equations of motion for a satellite under perturbations are:

$$\ddot{\mathbf{r}} = -\frac{\mu}{r^3}\mathbf{r} + \mathbf{a}_p \quad (49)$$

where $\mathbf{a}_p$ is the sum of all perturbing accelerations. Two main approaches exist for numerical integration:

12.1 Cowell’s Method (Direct Integration)

Cowell’s method [3, 4, 5] directly integrates the total acceleration:

$$\dot{\mathbf{r}} = \mathbf{v} \quad (50)$$

$$\dot{\mathbf{v}} = -\frac{\mu}{r^3}\mathbf{r} + \mathbf{a}_p \quad (51)$$

Advantages:

- Simplicity of formulation and implementation
- Any number of perturbations can be handled simultaneously
- Obtains $\mathbf{r}$ and $\mathbf{v}$ directly

Disadvantages:

- Near large attracting bodies, smaller steps are required (affects time and roundoff error)
- Approximately $10\times$ slower than Encke’s method [3]
- Roundoff error accumulates rapidly in double-precision for interplanetary flights with small step sizes

12.2 Encke’s Method (Osculating Orbit)

Encke’s method [3, 5] integrates the *deviation* from a reference (osculating) two-body orbit rather than the total state. Let $\rho(t)$ be the osculating orbit position (the path the satellite would follow if perturbations vanished at epoch t_0). The deviation $\delta\mathbf{r} = \mathbf{r} - \rho$ satisfies:

$$\delta\ddot{\mathbf{r}} = \mathbf{a}_p + \frac{\mu}{\rho^3} \left[(1 - 2q)^{-3/2} \mathbf{r} - \delta\mathbf{r} \right] \quad (52)$$

where $q = \delta\mathbf{r} \cdot (2\mathbf{r} - \delta\mathbf{r})/r^2$.

Advantages:

- Faster than Cowell’s method (deviations change slowly)
- Better accuracy for small perturbations
- Less accumulation of roundoff error

Disadvantages:

- More complex formulation
- Requires *rectification* (resetting the osculating orbit) when $\delta r/r$ exceeds a tolerance
- The term $(1 - 2q)^{-3/2}$ requires careful numerical handling (difference of nearly equal quantities)

12.3 Störmer/Numerov vs. Cowell for Orbital Mechanics

Our `StörmerStep` and `NumerovStep` implementations are **integration methods**, while Cowell’s and Encke’s are **formulation frameworks**. The relationship:

Formulation	Integration Method	Use Case
Cowell’s	Störmer (order 2) or Numerov (order 4)	Direct integration of $y'' = f(x, y)$
Cowell’s	RK4, Adams-Moulton	General first-order reduction
Encke’s	Störmer/Numerov	Deviation from osculating orbit

Why Numerov outperforms standard Cowell integration:

1. **Order 4 accuracy** vs. order 2 for explicit Störmer or standard RK4 on first-order form
2. **Same stencil** (3 points) as explicit Störmer, but implicit correction achieves higher order
3. **Symplectic-like** energy behavior for near-Keplerian orbits (nearly circular)
4. **Fewer function evaluations** than equivalent RK4: 3 f -evals/step vs. 4 for RK4

Bate et al. [3] note that for orbital problems, the **Gauss-Jackson** (sum-squared, Σ^2) method is "one of the best, and most used, numerical integration methods for trajectory problems of the Cowell and Encke type." The Gauss-Jackson method is also a multistep predictor-corrector for second-order equations, similar in spirit to Numerov but with higher order and specialized for orbital mechanics.

12.4 Common Perturbations in Astrodynamics

Based on Curtis [4], Hintz [5], and Bate [3], the primary perturbations affecting Earth satellites are:

1. **Earth Oblateness (J_2):** The Earth's equatorial bulge causes secular variations in Ω (RAAN) and ω (argument of perigee). The nodal regression rate is:

$$\dot{\Omega} = -\frac{3J_2}{2} \sqrt{\frac{\mu}{a^3}} \left(\frac{R_\oplus}{a(1-e^2)} \right)^2 \cos i \quad (53)$$

2. **Atmospheric Drag:** Significant for LEO satellites. The acceleration is:

$$\mathbf{a}_d = -\frac{1}{2} \rho v_{rel}^2 \frac{C_D A}{m} \frac{\mathbf{v}_{rel}}{|\mathbf{v}_{rel}|} \quad (54)$$

where ρ is atmospheric density, $C_D \approx 2.0$ – 2.2 for typical satellites.

3. **Third-Body Gravity (Moon, Sun):** The lunar perturbation acceleration on an Earth satellite is:

$$\mathbf{a}_p = \mu_M \left(\frac{\mathbf{r}_{ms}}{r_{ms}^3} - \frac{\mathbf{r}_{Me}}{r_{Me}^3} \right) \quad (55)$$

where $\mathbf{r}_{ms}$ is the vector from Moon to satellite, $\mathbf{r}_{Me}$ from Moon to Earth.

4. **Solar Radiation Pressure:** For spacecraft with large area-to-mass ratio:

$$\mathbf{a}_{SRP} = -p_{SR} \frac{A}{m} \cos \theta \hat{\mathbf{s}} \quad (56)$$

where $p_{SR} \approx 4.5 \times 10^{-6}$ N/m² at 1 AU.

Special vs. General Perturbations:

- **Special perturbations:** Direct numerical integration (Cowell, Encke) — computes specific trajectories for given initial conditions.
- **General perturbations:** Analytic integration of series expansions (variation of parameters, Lagrange planetary equations) — provides analytical insight but is more complex.

12.5 Application: Numerov Orbit Propagation Demo

The `NumerovOrbitDemo.cpp` implementation demonstrates two scenarios:

Toy Example: Pure Two-Body LEO (400 km, 28.5° inclination)

- **Purpose:** Verify Numerov implementation correctness
- **Physics:** Central gravity only ($\mathbf{a} = -\mu\mathbf{r}/r^3$)
- **Validation:** Energy conservation to $\sim 0.04\%$ over 10 orbits
- **Expected:** Perfectly periodic orbit with constant energy

Real-World Example: LEO with J_2 Perturbation

- **Purpose:** Demonstrate dominant LEO perturbation effect
- **Physics:** Two-body + Earth oblateness ($J_2 = 1.08263 \times 10^{-3}$)
- **Perturbation formula** (Curtis §10.5):

$$\mathbf{a}_{J_2} = -\frac{3\mu J_2 R_\oplus^2}{2r^5} \begin{bmatrix} x(1 - 5z^2/r^2) \\ y(1 - 5z^2/r^2) \\ z(3 - 5z^2/r^2) \end{bmatrix} \quad (57)$$

- **Observable effects:**
 - Nodal regression: $-7.08/\text{day}$ (Curtis Eq. 4.52)
 - Perigee advance: $+11.52/\text{day}$ (Curtis Eq. 4.53)
- **Why nontrivial:** J_2 is the dominant LEO perturbation ($\sim 90\%$ of all perturbation effects)

Build and run (from Cosmos/Source):

```
mkdir -p build && cd build
cmake .. && make NumerovOrbitDemo
./Examples/NumerovOrbitDemo
```

Output: `numerov_twobody.csv`, `numerov_j2.csv` with columns `t,orbit,r_x,r_y,r_z,energy,error`.

13 Summary

1. Störmer’s rule exploits $y'' = f(x, y)$ (no y' in f) via the three-point recurrence $y_{m+1} = 2y_m - y_{m-1} + h^2 f_m$.
2. In δ -form: $\delta_m = \delta_{m-1} + h^2 f_m$, with starting value $\delta_0 = hv_0 + \frac{h^2}{2} f_0$. This form is mandatory to avoid numerical instability from the double root $\varrho(\zeta) = (\zeta - 1)^2$ (Section 3.3).
3. The algorithm is equivalent to Störmer-Verlet (leapfrog), which is symplectic for Hamiltonian systems.
4. “Order p ” means global error $\mathcal{O}(h^p)$ in step size h (not polynomial degree). The explicit method is order 2; Numerov’s implicit method (40) saturates the order barrier for the two-step stencil at order 4.
5. Gragg’s theorem guarantees an even-power error series in h , enabling efficient Bulirsch-Stoer Richardson extrapolation.
6. NR3’s `StepperStoerm` implements the $k = 2$ explicit Störmer in one-step extrapolation mode (HNW §II.14, Eq. 14.32) — the simplest special case of the general multistep family (HNW §III.10).
7. Numerov PECECE (Section 11) achieves order 4 on the same three-point stencil at 3 f -evaluations/step vs. 1 for explicit Störmer. Implemented as `NumerovStep`.

References.

References

- [1] E. Hairer, S. P. Nørsett, G. Wanner, *Solving Ordinary Differential Equations I: Nonstiff Problems*, 2nd ed. (Springer, 1993). Key sections: §II.14 (Nyström methods, extrapolation for $y'' = f(x, y)$), §III.10 (multistep Störmer methods, pp. 461–473).
- [2] W. B. Gragg, “On extrapolation algorithms for ordinary initial value problems,” *SIAM J. Numer. Anal.* **2** (1965) 384–403.
- [3] R. R. Bate, D. D. Mueller, J. E. White, *Fundamentals of Astrodynamics* (Dover, 1971). §9.2 (Cowell’s Method), §9.3 (Encke’s Method), §9.6 (Numerical Integration Methods including Gauss-Jackson).
- [4] H. D. Curtis, *Orbital Mechanics for Engineering Students*, 2nd ed. (Elsevier, 2005). Ch. 10 (Orbital Perturbations): §10.2 (Cowell), §10.3 (Encke), §10.4 (Atmospheric Drag), §10.5 (Gravitational Perturbations).
- [5] G. R. Hintz, *Orbital Mechanics and Astrodynamics* (Springer, 2015). Ch. 5: §5.2 (Cowell’s Method), §5.3 (Encke’s Method), §5.5 (Integration Schemes), §5.7 (Analytic Formulation of Perturbations).